

Assignment 1: Linux Commands & Shell Programming

1. What is a shell?

→ A shell is a command-line interface between the user and the kernel. It interprets user commands and executes them in the operating system.

2. What is a shell script?

→ A shell script is a text file containing a sequence of shell commands that are executed together like a program to automate tasks.

3. What is exit status in Linux?

→ Every command returns an exit code: 0 indicates successful execution, while any non-zero value indicates an error or failure.

4. What does \$? represent?

→ It stores the exit status of the last executed command and is useful for checking whether the previous command succeeded.

5. Differentiate user-defined and system variables.

→ System variables (like \$HOME, \$PATH) are predefined by Linux, whereas user-defined variables are created by users for their own scripts.

6. Explain the purpose of the test command.

→ test (or []) checks conditions such as file existence, numeric or string comparisons, and returns true or false status.

7. What loops are supported in shell scripting?

→ Shell supports for, while, and until loops for repeated execution of commands.

8. What does the grep command do?

→ It searches for specific patterns or strings within files and displays matching lines.

9. What is the function of sed?

→ sed (stream editor) performs text editing tasks like search, replace, insertion, or deletion without opening the file.

10. What does the man command do?

→ It displays the manual page for any Linux command, explaining its syntax, options, and usage examples.

Assignment 2: Process Control System Calls

- 1. Define a process.**
→ A process is an executing instance of a program, having its own address space, resources, and process ID (PID).
- 2. What is fork() used for?**
→ fork() creates a new process called a child, which runs concurrently with the parent process.
- 3. What does fork() return?**
→ It returns 0 to the child, the child's PID to the parent, and -1 on error.
- 4. Explain the use of exec() family of calls.**
→ These replace the current process image with a new program, allowing one process to run another executable.
- 5. What is the purpose of wait()?**
→ It makes the parent process pause until its child finishes, ensuring proper synchronization and cleanup.
- 6. What is a zombie process?**
→ A process that has terminated but whose parent hasn't collected its exit status yet; it remains in the process table temporarily.
- 7. What is an orphan process?**
→ A child process whose parent has terminated; it is automatically adopted by the init process.
- 8. What does getpid() return?**
→ The unique process ID of the calling process.
- 9. What does getppid() return?**
→ The PID of the parent process.
- 10. What is system() function used for?**
→ It executes a shell command directly from a C program, similar to typing it in a terminal.

Assignment 3: CPU Scheduling (SJF & Round Robin)

- 1. Define burst time.**
→ The CPU time required by a process to complete execution once it gets the CPU.

2. **What is waiting time?**
→ The total time a process spends in the ready queue waiting to be executed.
3. **What is turnaround time?**
→ The total time from process submission to its completion, i.e., waiting + burst time.
4. **Explain Shortest Job First (SJF).**
→ In SJF, the process with the smallest burst time executes first, giving minimal average waiting time.
5. **Differentiate preemptive and non-preemptive scheduling.**
→ In preemptive, a process can be interrupted; in non-preemptive, a process runs until completion.
6. **What is Round Robin scheduling?**
→ Each process gets a fixed time quantum to run in cyclic order; suitable for time-sharing systems.
7. **What is time quantum or time slice?**
→ The small fixed CPU time allotted to each process in Round Robin scheduling.
8. **Which algorithm gives minimum average waiting time?**
→ Shortest Job First (SJF).
9. **List CPU scheduling criteria.**
→ CPU utilization, throughput, turnaround time, waiting time, and response time.
10. **What is context switching?**
→ The act of saving the current process state and loading another process's state when CPU changes.



Assignment 4A: Thread Synchronization (Producer–Consumer)

1. **What is a semaphore?**
→ A semaphore is a synchronization tool used to control access to shared resources by multiple processes.
2. **Types of semaphores?**
→ Binary (0/1) for mutual exclusion and Counting for controlling multiple resource instances.

3. **Explain wait() and signal() operations.**
→ wait() decrements the semaphore and may block if zero; signal() increments it and unblocks waiting processes.
 4. **What is a mutex?**
→ A mutual-exclusion lock that allows only one thread to access a shared resource at a time.
 5. **What is the critical section problem?**
→ Ensuring that only one process executes a critical section (shared data access) at any given time.
 6. **Explain the producer-consumer problem.**
→ It involves two processes sharing a buffer — the producer adds items, and the consumer removes them using semaphores for synchronization.
 7. **What happens if semaphore value < 0?**
→ Processes performing wait() are blocked until another process performs signal().
 8. **What is the initial value of a mutex?**
→ Usually 1, representing an unlocked state.
 9. **List common pthread mutex functions.**
→ pthread_mutex_init, pthread_mutex_lock, pthread_mutex_unlock, pthread_mutex_destroy.
10. **What is deadlock?**
→ A situation where two or more processes wait indefinitely for resources held by each other.

Assignment 4B: Reader–Writer Problem

1. **What is the Reader–Writer problem?**
→ It deals with synchronization between processes that read and write a shared data area, ensuring consistency.
2. **Can multiple readers read simultaneously?**
→ Yes, multiple readers can read at the same time since they don't modify data.
3. **Can a writer write while readers are reading?**
→ No, writers must wait until all readers finish to avoid inconsistent data.

4. What is writer starvation?

→ When continuous reader activity prevents writers from ever accessing the shared resource.

5. Which synchronization tools are used?

→ Semaphores and mutexes are used to coordinate reader and writer access.

6. Which semaphore ensures mutual exclusion for writers?

→ The wsem semaphore ensures that only one writer can write at a time.

7. What does reader priority mean?

→ Readers are given access first if other readers are already reading, delaying writers.

8. What is the role of the readcount variable?

→ It tracks how many readers are currently reading; first reader locks the resource, last reader releases it.

9. Which header file provides thread synchronization functions?

→ pthread.h.

10. What is the initial value of wsem?

→ It is initialized to 1, meaning the shared data is initially free for access.



Assignment 5: Deadlock Avoidance – Banker's Algorithm

1. What is a deadlock?

→ A situation where two or more processes wait indefinitely for resources held by each other.

2. What is the Banker's algorithm used for?

→ It is used to avoid deadlock by checking whether a system is in a safe or unsafe state before allocating resources.

3. What is a safe state?

→ A state where all processes can complete their execution without leading to a deadlock.

4. What is an unsafe state?

→ A state where processes may not complete, potentially leading to deadlock.

5. What are the matrices used in the Banker's algorithm?

→ Allocation, Max, Need, and Available matrices.

6. **How is the Need matrix calculated?**
→ $\text{Need}[i][j] = \text{Max}[i][j] - \text{Allocation}[i][j]$
 7. **What is the condition for resource allocation in Banker's algorithm?**
→ $\text{Request} \leq \text{Need}$ and $\text{Request} \leq \text{Available}$.
 8. **What does the algorithm do after a resource request?**
→ Temporarily allocates resources, checks for a safe sequence, and grants only if safe.
 9. **Why is it called the Banker's algorithm?**
→ Because it models how a banker safely allocates loans to customers ensuring all can be repaid.
 10. **What happens if the system is unsafe?**
→ The resource request is denied to avoid potential deadlock.
-



Assignment 6: Page Replacement Algorithms (FCFS, LRU, Optimal)

1. **What is page replacement?**
→ It's the process of swapping pages between main memory and secondary storage when memory is full.
2. **What causes a page fault?**
→ When a process tries to access a page not currently in memory.
3. **Explain FCFS page replacement.**
→ The oldest loaded page is replaced first, regardless of how frequently it's used.
4. **Explain LRU (Least Recently Used).**
→ The page that has not been used for the longest time is replaced.
5. **Explain the Optimal page replacement algorithm.**
→ Replaces the page that will not be used for the longest time in the future.
6. **Which page replacement algorithm gives the least number of page faults?**
→ The Optimal algorithm.
7. **What is Belady's anomaly?**
→ In some cases, increasing frame size increases page faults (seen in FCFS).

8. **How does LRU approximate the Optimal algorithm?**
→ By using past behavior (recent use) to predict future use.
 9. **What is a reference string?**
→ A sequence of memory page numbers representing the order of page requests.
 10. **Why is the Optimal algorithm not practical?**
→ Because future page references cannot be predicted in real systems.
-

Assignment 7: Inter-Process Communication (FIFO & Shared Memory)

(A) FIFO Communication

1. **What is IPC?**
→ Inter-Process Communication allows processes to exchange data or signals.
 2. **What is FIFO?**
→ A named pipe that allows communication between unrelated processes in a first-in, first-out manner.
 3. **How does FIFO differ from an unnamed pipe?**
→ FIFO exists as a special file in the filesystem, unlike unnamed pipes which exist only during process execution.
 4. **Which system call is used to create a FIFO?**
→ mkfifo()
 5. **Can FIFO provide full-duplex communication?**
→ Yes, two FIFOs can be used for full-duplex (two-way) communication.
-

(B) Shared Memory Communication

6. **What is shared memory?**
→ A memory segment accessible by multiple processes for high-speed data exchange.
7. **Which system calls are used for shared memory?**
→ shmget(), shmat(), shmdt(), and shmctl().

8. What is the advantage of shared memory IPC?

→ It is the fastest form of IPC because processes directly access the same memory.

9. What synchronization mechanism is used with shared memory?

→ Semaphores or mutex locks to prevent concurrent access issues.

10. What happens if synchronization is not maintained?

→ Data inconsistency or race conditions can occur.



Assignment 8: Disk Scheduling Algorithms (SSTF, SCAN, C-LOOK)

1. What is disk scheduling?

→ It determines the order in which I/O requests to the disk are serviced to improve performance.

2. What is SSTF (Shortest Seek Time First)?

→ The disk head moves to the request closest to its current position.

3. What is SCAN (Elevator) algorithm?

→ The disk arm moves in one direction, servicing requests, then reverses direction like an elevator.

4. What is C-LOOK?

→ Similar to SCAN, but the head moves in one direction and jumps back to the start without servicing on return.

5. Which disk scheduling algorithm gives minimum average seek time?

→ SSTF usually gives better average seek time than FCFS.

6. What is seek time?

→ The time taken by the disk arm to move to the track containing the desired sector.

7. What is rotational latency?

→ The delay waiting for the desired disk sector to rotate under the read-write head.

8. Which algorithm prevents starvation?

→ SCAN and C-LOOK reduce starvation compared to SSTF.

9. What does “initial head position” mean?

→ The starting track number of the disk arm before servicing begins.

10.What is the main goal of disk scheduling?

→ To minimize seek time and improve throughput of I/O operations.

❖ Assignment 9: Study Assignment – System Call Implementation

1. What is a system call?

→ A request from user space to the kernel to perform a privileged operation, like I/O or process management.

2. How are system calls implemented?

→ By adding a function in kernel source, registering it in the system call table, recompiling the kernel, and calling it from user space.

3. Example of a system call in Linux?

→ fork(), read(), write(), open(), close(), exec().

4. What is the difference between user mode and kernel mode?

→ User mode restricts direct hardware access, while kernel mode has full control over hardware and memory.

5. Why are system calls needed?

→ They provide controlled access to critical kernel resources securely.

6. What happens when a system call is invoked?

→ CPU switches from user mode to kernel mode to execute the kernel-defined routine.

7. How can you identify system calls in C code?

→ By using commands like strace to trace system calls made by a program.

8. What is kernel compilation?

→ The process of building a custom Linux kernel after modifying its source code.

9. What are examples of system call categories?

→ Process control, file management, device management, and inter-process communication.

10.What is the purpose of adding a custom system call?

→ To extend kernel functionality or add custom features for experimentation or research.